

**Analysis Document**  
CS 6012  
Lindsay Haslam

**Who are your team members?**

Corinne Jones.


**Mergesort Threshold Experiment**

In the following experiment, we discovered that the optimal threshold value for the given array predominantly lies around 20. While there were occasional variations,, with a threshold of 10 outperforming in some instances, the consistently superior performance was with a threshold value of 20. Also, the threshold corresponding to the maximum value (denoted in red) does not resort to insertion sort. The red line increases in execution time when in contrast with alternative threshold values.

Threshold 5		Threshold 20	
100	329458	100	12709
500	402083	500	102417
1000	1240916	1000	203666
5000	2243542	5000	1410125

Threshold 10		Threshold 50	
100	30959	100	20417
500	296791	500	104583
1000	192625	1000	241208
5000	1390583	5000	1582625

ThreshMaxVal	
100	8916
500	202375
1000	793417
5000	19968791



## Quicksort Pivot

The best pivot point identified for quicksort was the median index. This choice aligns logically with our approach, where we consider the first and last indices (represented by 0 and `array.length-1`, respectively) and calculate their midpoint. Professor Jones, in his lectures, asserted that the median index consistently emerges as the optimal choice, with random selection ranking at second.

Random Index


10	76542
100	182083
500	440791
1000	320666
5000	1916000
7000	796542
10000	859625

Left Index

10	205875
100	233208
500	516542
1000	290875
5000	1871792
7000	2628458
10000	965375

Median Index

10	13042
100	38542
500	79416
1000	171917
5000	1148917
7000	169833
10000	212583



## Quick Sort vs Merge Sort


After running best, average, and worst case. It is apparent that quicksort is the ideal sort every time. When it comes to average sort, the differences between the two aren't as drastic, as they seem to follow a similar pattern. For best case, merge sort takes a significant amount of time for smaller sized arrays, and drops at a certain size.

Merge

10	50083
100	136959
500	3953250
1000	976458
5000	1608125
7000	1159834
10000	2494625

Quick

10	19500
100	183542
500	583000
1000	199959
5000	711459
7000	934916
10000	1077958

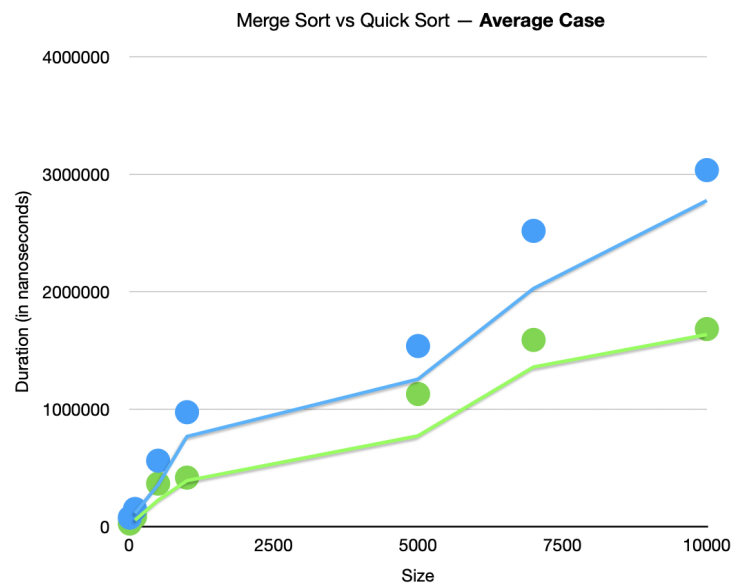


Merge

10	77458
100	150875
500	561500
1000	974084
5000	1536958
7000	2516459
10000	3033292

Quick

10	27000
100	84375
500	366625
1000	417417
5000	1128666
7000	1588833
10000	1681417




Merge

10	71083
100	310959
500	565541
1000	897750
5000	2064583
7000	1780292
10000	2787250

Quick

10	21792
100	124208
500	254916
1000	193459
5000	589083
7000	801708
10000	1636208



**Do the actual running times of your sorting methods exhibit the growth rates you expected to see?**

I did not anticipate that different thresholds would create such an impact on the efficiency of merge and insertion sort. It was genuinely surprising to see the substantial difference made by entering insertion sort either too early or too late in size. I'm interested to see how engineers frequently combine various algorithms to achieve optimal runtimes.

I was expecting the median index to be the best pivot, but I was surprised that the random index method's results weren't closer to the median's results. I think a better way to test the random index method would be to iterate through it many times (meaning a new random array is created each time) and create some sort of average with it. The specific graph that I have in this document is a representation of one array, and it could have been a record of any index. There's a chance the results would have been better than, closer to, or drastically different than median index's results.

As for the growth rates in best, average, and worst cases, my expectations were kind of met. Professor Jones mentioned that merge sort tends to be slower and less efficient, so I assumed that at least one case might favor merge sort. While I initially expected most graphs to lean towards quicksort, I was surprised that *all* scenarios consistently demonstrate quicker durations with quicksort. It was also interesting to see how sometimes my lines didn't look "obviously" like the Big O notation that it should.. Occasionally, by zooming out from the graph or using logarithmic settings in Numbers, the resemblance to a specific Big O notation becomes clearer.